

NAME

`irsim` – An event-driven logic-level simulator for MOS circuits

SYNOPSIS

irsim [*-s*] *prm_file sim_file ...* [*+hist_file*] [*-cmd_file ...*]

DESCRIPTION

IRSIM is an event-driven logic-level simulator for MOS (both N and P) transistor circuits. Two simulation models are available:

switch

Each transistor is modeled as a voltage-controlled switch. Useful for initializing or determining the functionality of the network.

linear

Each transistor is modeled as a resistor in series with a voltage-controlled switch; each node has a capacitance. Node values and transition times are computed from the resulting RC network, using Chorng-Yeoung Chu's model. Chris Terman's original model is not supported any more.

If the **-s** switch is specified, 2 or more transistors of the same type connected in series, with no other connections to their common source/drain will be *stacked* into a compound transistor with multiple gates.

The **prm_file** is the electrical parameters file that configure the devices to be simulated. It defines the capacitance of the various layers, transistor resistances, threshold voltages, etc... (see `presim(1)`). If *prm_file* does not specify an absolute path then IRSIM will search for the *prm_file* as follows (in that order):

- 1) `./<prm_file>` (in the current directory).
- 2) `~cad/lib/<prm_file>`
- 3) `~/cad/lib/<prm_file>.prm`

If `~cad/` does not exist, IRSIM will try to use directory `/projects/cad`. The default search directory (`~cad`) can be overridden by setting the environment variable `CAD_HOME` to the appropriate directory prior to running IRSIM (i.e. `setenv CAD_HOME /usr/beta/mycad`).

IRSIM first processes the files named on the command line, then (assuming the exit command has not been processed) accepts commands from the user, executing each command before reading the next.

File names NOT beginning with a `'-'` are assumed to be sim files (see `sim(5)`), note that this version does not require to run the sim files through `presim`. These files are read and added to the network database. There is only a single name space for nodes, so references to node "A" in different network files all refer to the same node. While this feature allows one to modularize a large circuit into several network files, care must be taken to ensure that no unwanted node merges happen due to an unfortunate clash in names.

File names prefaced with a `'-'` are assumed to be command files: text files which contain command lines to be processed in the normal fashion. These files are processed line by line; when an end-of-file is encountered, processing continues with the next file. After all the command files have been processed, and if an "exit" command has not terminated the simulation run, IRSIM will accept further commands from the user, prompting for each one like so:

irsim>

The **hist_file** is the name of a file created with the *dumph* command (see below). If it is present, IRSIM will initialize the network to the state saved in that file. This file is different from the ones created with the `">"` command since it saves the state of every node for all times, including any pending events.

This version supports changes to the network through the **update** command. Also, the capability to incrementally re-simulate the network up to the current time is provided by the **isim** command.

COMMAND SUMMARY

@ filename take commands from command file
? wnode... print info about node's source/drain connections
! wnode... print info about node's gate connections
< filename restore network state from file
> filename write current network state to file
<< filename same as "<" but restores inputs too
| comment... comment line
activity from [to] graph circuit activity in time interval
ana wnode... display nodes in analyzer window
analyzer wnode... display nodes in analyzer window
assert wnode [m] val assert that *wnode* equals *value*
back [time] move back to *time*
c [n] simulate for *n* clock cycles (default:1)
changes from [to] print nodes that changed in time interval
clock [node [val]] define value sequence for clock node
clear clear analyzer window (remove signals)
d [wnode]... print display list or specified node(s)
debug [debug_level...]
 set debug level (default: off)
decay [n] set charge decay time (0 => no decay)
display [arg]... control what gets displayed when
dumph filename... write net history to file
exit [status] return to system
flush [time] flush out history up to *time* (default: now)
h wnode... make node logic high (1) input
has_coords print YES if transistor coordinates are available
inputs print current list of input nodes
ires [n] set incremental resolution to *n* ns
isim [filename] incrementally resimulate changes from *filename*
l wnode... make node logic low (0) input
logfile [filename] start/stop log file
model [name] set simulation model to *name*
p step clock one simulation step (phase)
path wnode... display critical path for last transition of a node
print comment... print specified text
printp print a list of all pending events
printx print all undefined (X) nodes
q terminate input from current stream
R [n] simulate for *n* cycles (default:longest sequence)
readh filename read history from *filename*
report[level] set/reset reporting of decay events
s [n] simulate for *n* ns. (default: stepsize)
stepsize [n] set simulation step size to *n* ns.
set vector value assign *value* to *vector*
setlog[file/off] log net changes to file (*off* -> no log)
"setpath" "[path...]"
 set search path for cmd files
stats print event statistics
t [-]wnode... start/stop tracing of specified nodes
tcap print list of shorted transistors
"time [command] print resource utilization summary
u wnode... make node undefined (X) input
unitdelay [n] force transitions to take *n* ns. (0 disables)
update filename read net changes from file
V [node [value...]] define sequence of inputs for a node
vector label node... define bit vector
w [-]wnode... add/delete nodes from display list
wnet [filename] write network to file
x wnode... remove node from input lists
Xdisplay[host:n] set/show X display (for analyzer)

COMMAND DESCRIPTIONS

Commands have the following simple syntax:

cmd *arg1 arg2 ... argn* <newline>

where **cmd** specifies the command to be performed and the *argi* are arguments to that command. The arguments are separated by spaces (or tabs) and the command is terminated by a <newline>.

If **cmd** is not one of the built-in commands documented below, IRSIM appends ".cmd" to the command name and tries to open that file as a command file (see "@" command). Thus the command "foo" has the same effect as "@ foo.cmd".

Notation:

...
indicates zero or more repetitions

[] enclosed arguments are optional

node name of node or vector in network

wnode name of node or vector in network, can include '*' wildcard which matches any sequence of zero or more characters. The pair of characters '{' and '}' denote iteration over the limits enclosed by it, for example: **name{1:10}** will expand into *name1, name2 ... name10*. A 3rd optional argument sets the stride, for example: **name{1:10:2}** will expand into *name1, name3, ... name7, name9*.

| **comment...**

Lines beginning with vertical bar are treated as comments and ignored -- useful for comments or temporarily disabling certain commands in a command file.

Most commands take one or more node names as arguments. Whenever a node name is acceptable in a command line, one can also use the name of a bit vector. In this case, the command will be applied to each node of the vector (the "t" and "d" treat vectors specially, see below).

vector *label node...*

Define a bit vector named "label" which includes the specified nodes. If you redefine a bit vector, any special attributes of the old vector (e.g., being on the display or trace list) are lost. Wild cards are not accepted in the list of node names since you would have no control over the order in which matching nodes would appear in the vector.

The simulator performs most commands silently. To find out what's happened you can use one of the following commands to examine the state of the network and/or the simulator.

set *vector value*

Assign value to vector. For example, the following sequence of commands:

vector BUS bit.1 bit.2 bit.3 **set** BUS 01x

The first command will define *BUS* to be a vector composed of nodes *bit.1*, *bit.2*, and *bit.3*. The second command will assign the following values:

bit.1 = 0
bit.2 = 1
bit.3 = X

Value can be any sequence of [0,1,h,H,l,L,x,X], and must be of the same length as the bit vector itself.

d [*wnode*]...

Display. Without arguments displays the values all nodes and bit vectors currently on the display list (see **w** command). With arguments, only displays the nodes or bit vectors specified. See also the "display" command if you wish to have the display list printed out automatically at the end of certain simulation commands.

w [-]*wnode*...

Watch/unwatch one or more nodes. Whenever a "d" command is given, each watched node will displayed like so:

node1=0 node2=X ...

To remove a node from the watched list, preface its name with a '-'. If *wnode* is the name of a bit vector, the values of the nodes which make up the vector will be displayed as follows:

label=010100

where the first 0 is the value of first node in the list, the first 1 the value of the second node, etc.

assert *wnode* [*mask*] *value*

Assert that the boolean value of the node or vector *wnode* *is value*. If the comparison fails, an error message is printed. If *mask* is given then only those bits corresponding to zero bits in *mask* take part in the comparison, any character other than 0 will skip that bit. The format of the error message is the following:

(tty, 3): assertion failed on 'name' 10X10 (1010X)

Where *name* is the name of the vector, followed by the actual value and the expected value enclosed in parenthesis. If a *mask* is specified, then bits that were not compared are printed as '-'.

ana *wnode*...

This is a shorthand for the analyzer command (described below).

analyzer *wnode*...

Add the specified node(s)/vector(s) to the analyzer display list (see irsim-analyzer(3) for a detailed explanation). If the analyzer window does not exist, it will be created. If no arguments are given and the analyzer window already exists, nothing happens.

Xdisplay [*host:display*]

You must be able to connect to an X-server to start the analyzer. If you haven't set up the DISPLAY environment variable properly, the analyzer command may fail. If this is the case you can use the Xdisplay command to set it from within the simulator. With no arguments, the name of the current X-server will be printed.

clear Removes all nodes and vectors from the analyzer window. This command is most useful in command scripts for switching between different signals being displayed on the analyzer.

"?" and "!" allow the user to go both backwards and forwards through the network. This is a useful debugging aid.

? *wnode*...

Prints a synopsis of the named nodes including their current values and the state of all transistors that affect the value of these nodes. This is the most common way of wandering through the network in search of what went wrong. The output from the command ? out looks like

```
out=0 (vl=0.3 vh=0.8) (0.100 pf) is computed from:
n-channel phi2=0 out=0 in=0 [1.0e+04, 1.3e+04, 8.7e+03]
pulled down by (a=1 b=1) [1.0e+04, 1.3e+04, 8.8e+03]
pulled up [4.0e+04, 7.4e+04, 4.0e+04]
```

The first line gives the node's name and current value, its low and high logic thresholds, user-specified low-to-high and high-to-low propagation delays if present, and its capacitance if nonzero.

Succeeding lines list the transistor whose sources or drains connect to this node: the transistor type ("pulled down" is an n-channel transistor connected to gnd, "pulled up" is a depletion pullup or p-channel transistor connected to vdd), the values of the gate, source, and drain nodes, and the modeling resistances. Simple chains of transistors with the same implant type are collapsed by the `-s` option into a single transistor with a "compound" gate; compound gates appear as a parenthesized list of nodes (e.g., the pulldown shown above). The three resistance values -- static, dynamic high, dynamic low -- are given in Kilo-ohms.

Finally, any pending events for a node are listed after the electrical information.

! wnode...

For each node in the argument list, print a list of transistors controlled by that node.

tcap Prints a list of all transistors with their source/drain shorted together or whose source/drain are connected to the power supplies. These transistors will have no effect on the simulation other than their gate capacitance load. Although transistors connected across the power supplies are real design errors, the simulator does not complain about them.

Any node can be made an input -- the simulator will not change an input node's value until it is released. Usually on specific nodes -- inputs to the circuit -- are manipulated using the commands below, but you can fool with a subcircuit by forcing values on internal nodes just as easily.

h wnode...

Force each node on the argument list to be a high (1) input. Overrides previous input commands if necessary.

l wnode...

Like "h" except forces nodes to be a low (0) input.

u wnode...

Like "h" except forces nodes to be a undefined (X) input.

x wnode...

Removes nodes from whatever input list they happen to be on. The next simulation step will determine the correct node value from the surrounding circuit. This is the default state of most nodes. Note that this does not force nodes to have an "X" value -- it simply removes them from the input lists.

inputs

prints the high, low, and undefined input lists.

It is possible to define a sequence of values for a node, and then cycle the circuit as many times as necessary to input each value and simulate the network. A similar mechanism is used to define the sequence of values each clock node goes through during a single cycle.

Each value is a list of characters (with no intervening blanks) chosen from the following:

- 1, h, H logic high (1)
- 0, l, L logic low (0)
- u, U undefined (X)
- x, X remove node from input lists

Presumably the length of the character list is the same as the size of the node/vector to which it will be assigned. Blanks (spaces and tabs) are used to separate values in a sequence. The sequence is used one value at a time, left to right. If more values are needed than supplied by the sequence, IRSIM just restarts the sequence again.

V [node [value...]]

Define a vector of inputs for a node. After each cycle of an "R" command, the node is set to the next value specified in the sequence.

With no arguments, clears all input sequences (does not affect clock sequences however). With one argument, "node", clears any input sequences for that node/vector.

clock [*node* [*value...*]]

Define a phase of the clock. Each cycle, each node specified by a clock command must run through its respective values. For example,

```
clock phi1 1 0 0 0
clock phi2 0 0 1 0
```

defines a simple 4-phase clock using nodes *phi1* and *phi2*. Alternatively one could have issued the following commands:

```
vector clk phi1 phi2
clock clk 10 00 01 00
```

With no arguments, clears all clock sequences. With one argument, "node", clears any clock sequences for that node/vector.

After input values have been established, their effect can be propagated through the network with the following commands. The basic simulated time unit is 0.1ns; all event times are quantized into basic time units. A simulation step continues until *stepsize* ns. have elapsed, and any events scheduled for that interval are processed. It is possible to build circuits which oscillate -- if the period of oscillation is zero, the simulation command will not return. If this seems to be the case, you can hit <ctrl-C> to return to the command interpreter. Note that if you do this while input is being taken from a file, the simulator will bring you to the top level interpreter, aborting all pending input from any command files.

When using the linear model (see the "**model**" command) transition times are estimated using an RC time constant calculated from the surrounding circuit. When using the switch model, transitions are scheduled with unit delay. These calculations can be overridden for a node by setting its *tphl* and *tphl* parameters which will then be used to determine the time for a transition.

s [*n*] **Simulation step. Propagates new values for the inputs through the network, returns when *n* (default: *stepsize*) ns. have passed.** If *n* is specified, it will temporarily override the *stepsize* value. Unlike previous versions, this value is NOT remembered as the default value for the *stepsize* parameter. If the display mode is "automatic", the current display list is printed out on the completion of this command (see "display" command).

c [*n*] **Cycle *n* times (default: 1) through the clock, as defined by the "**clock**" command. Each phase of the clock lasts *stepsize* ns. If the display mode is "automatic", the current display list is printed out on the completion of this command (see "**display**" command).**

p **Step the clock through one phase (or simulation step).** For example, if the clock is defined as above

```
clock phi1 1 0 0 0
clock phi2 0 0 1 0
```

then "**p**" will set *phi1* to 1 and *phi2* to 0, and then propagate the effects for one simulation step. The next time "**p**" is issued, *phi1* and *phi2* will both be set to 0, and the effects propagated, and so on. If the "**c**" command is issued after "**p**" has been used, the effect will be to step through the next 4 phases from where the "**p**" command left off.

R [*n*] **Run the simulator through *n* cycles (see the "**c**" command). If *n* is not present make the run as long as the longest sequence. If display mode is automatic (see "**display**" command) the display is printed at the end of each cycle. Each "**R**" command starts over at the beginning of the sequence defined for each node.**

back time

Move back to the specified time. This command restores circuit state as of *time*, effectively undoing any changes in between. Note that you can not move past any previously flushed out history (see flush command below) as the history mechanism is used to restore the network state. This command can be useful to undo a mistake in the input vectors or to re-simulate the circuit with a different debug level.

path *wnode...*

display critical path(s) for last transition of the specified node(s). The critical path transistions are reported using the following format:

node -> value @ time (delta)

where *node* is the name of the node, *value* is the value to which the node transitioned, *time* is the time at which the transistion occurred, and *delta* is the delay through the node since the last transition. For example:

```
critical path for last transition of Hit_v1:
phi1-> 1 @ 2900.0ns , node was an input
PC_driver-> 0 @ 2900.4ns (0.4ns)
PC_b_q1-> 1 @ 2904.0ns (3.6ns)
tagDone_b_v1-> 0 @ 2912.8ns (8.8ns)
tagDone1_v1-> 1 @ 2915.3ns (2.5ns)
tagDone1_b_v1-> 0 @ 2916.0ns (0.7ns)
tagDone_v1-> 1 @ 2918.4ns (2.4ns)
tagCmp_b_v1-> 0 @ 2922.1ns (3.7ns)
tagCmp_v1-> 1 @ 2923.0ns (0.9ns)
Vbit_b_v1-> 0 @ 2923.2ns (0.2ns)
Hit_v1-> 1 @ 2923.5ns (0.3ns)
```

activity *from_time [to_time]*

print histogram showing amount of circuit activity in the specified time interval. Actually only shows number of nodes which had their most recent transition in the interval.

changes *from_time [to_time]*

print list of nodes which last changed value in the specified time interval.

printp

print list of all pending events sorted in time. The node associated with each event and the scheduled time is printed.

printx

print a list of all nodes with undefined (X) values.

Using the trace command, it is possible to get more detail about what's happening to a particular node. Much of what is said below is described in much more detail in "Logic-level Simulation for VLSI Circuits" by Chris Terman, available from Kluwer Academic Press. When a node is traced, the simulator reports each change in the node's value:

```
[event #100] node out.1: 0 -> 1 @ 407.6ns
```

The event index is incremented for each event that is processed. The transition is reported as

old value -> new value @ report time

Note that since the time the event is processed may differ from the event's report time, the report time for successive events may not be strictly increasing.

Depending on the debug level (see the "**debug**" command) each calculation of a traced node's value is reported:

```
[event #99] node clk: 0 -> 1 @ 400.2ns
final_value( Load ) V=[0.00, 0.04] => 0
..compute_tau( Load )
{Rmin=2.2K Rdom=2.2K Rmax=2.2K} {Ca=0.06 Cd=0.17}
tauA=0.1 tauD=0.4 ns
[event #99: clk->1] transition for Load: 1 -> 0 (tau=0.5ns, delay=0.6ns)
```

In this example, a calculation for node *Load* is reported. The calculation was caused by event 99 in which node *clk* went to 1. When using the linear model (as in this example) the report shows

current value -> final value

The second line displays information regarding the final value (or dc) analysis for node "Load"; the minimum and maximum voltages as well as the final logical value (0 in this case).

The next three lines display timing analysis information used to estimate the delays. The meaning of the variables displayed can be found Chu's thesis: "Improved Models for Switch-Level Simulation".

When the *final value* is reported as "D", the node is not connected to an input and may be scheduled to decay from its current value to X at some later time (see the "**decay**" command).

"tau" is the calculated transition time constant, "delta" is when any consequences of the event will be computed; the difference in the two times is how IRSIM accounts for the shape of the transition waveform on subsequent stages (see reference given above for more details). The middle lines of the report indicate the Thevenin and capacitance parameters of the surrounding networks, i.e., the parameters on which the transition calculations are based.

debug [*ev dc tau taup tw spk*][*off*][*all*]

Set debugging level. Useful for debugging simulator and/or circuit at various levels of the computation. The meaning of the various debug levels is as follows:

ev display event enqueueing and dequeuing.

dc display dc calculation information.

tau display time constant (timing) calculation.

taup display second time constant (timing) calculation.

tw display network parameters for each stage of the tree walk, this applies to **dc**, **tau**, and **taup**. This level of debugging detail is usually needed only when debugging the simulator.

spk displays spike analysis information.

all This is a shorthand for specifying all of the above.

off This turns off all debugging information.

If a debug switch is on then during a simulation step, each time a watched node is encountered in some event, that fact is indicated to the user along with some event info. If a node keeps appearing in this printout, chances are that its value is oscillating. Vice versa, if your circuit never settles (i.e., it oscillates), you can use the "**debug**" and "**t**" commands to find the node(s) that are causing the problem.

Without any arguments, the debug command prints the current debug level.

t [-]wnode...

set trace flag for node. Enables the various printouts described above. Prefacing the node name with '-' clear its trace flag. If "wnode" is the name of a vector, whenever any node of that vector changes value, the current time and the values of all traced vectors is printed. This feature is useful for watching the relative arrival times of values at nodes in an output vector.

System interface commands:

> *filename*

Write current state of each node into specified file. Useful for making a breakpoint in your simulation run. Only stores values so isn't really useful to "dump" a run for later use, i.e., the current input lists, pending events, etc. are NOT saved in the state file.

< *filename*

Read from specified file, reinitializing the value of each node as directed. Note that network

must already exist and be identical to the network used to create the dump file with the ">" command. These state saving commands are really provided so that complicated initializing sequences need only be simulated once.

<< *filename*

Same as "<" command, except that this command will restore the *input* status of the nodes as well. It does not, however, restore pending events.

dump *[filename]*

Write the history of the simulation to the specified file, that is; all transistions since time = 0. The resulting file is a machine-independent binary file, and contains all the required information to continue simulation at the time the dump takes place. If the filename isn't specified, it will be constructed by taking the name of the *sim_file* (from the command line) and appending ".hist" to it.

readh *filename*

Read the specified history-dump file into the current network. This command will restore the state of the circuit to that of the dump file, overwriting the current state.

flush *[time]*

If memory consumption due to history maintainance becomes prohibitive, this command can be used to free the memory consumed by the history up to the time specified. With no arguments, all history up to the current point in the simulation is freed. Flushing out the history may invalidate an incremental simulation and the portions flushed will no longer appear in the analyzer window.

setpath *[path...]*

Set the search-path for command files. *Path* should be a sequence of directories to be searched for ".cmd" files, "." meaning the current directory. For example:

setpath . /usr/me/rsim/cmds /cad/lib/cmds

With no arguments, it will print the current search-path. Initially this is just ".".

print *text...*

Simply prints the text on the user's console. Useful for keeping user posted of progress through a long command file.

logfile *[filename]*

Create a logfile with the specified name, closing current log file if any; if no argument, just close current logfile. All output which appears on user's console will also be placed in the logfile. Output to the logfile is cleverly formatted so that logfiles themselves can serve as command files.

setlog *[filename | off]*

Record all net changes, as well as resulting error messages, to the specified file (see "update" command). Net changes are always appended to the log-file, preceding each sequence of changes by the current date. If the argument is *off* then net-changes will not be logged. With no arguments, the name of the current log-file is printed.

The default is to always record net changes; if no filename is specified (using the "setlog" command) the default filename *irsim_changes.log* will be used. The log-files are formatted so that log-files may themselves be used as net-change files.

wnet *[filename]*

Write the current network to the specified file. If the filename isn't specified, it will be constructed by taking the name of the *sim_file* (from the command line) and appending ".inet" to it. The resulting file can be used in a future simulation run, as if it were a sim file. The file produced is a machine independent binary file, which is typically about 1/3 the size of the sim file and about 8 times faster to load.

time *[command]*

With no argument, a summary of time used by the simulator is printed. If arguments are given

the specified command is timed and a time summary is printed when the command completes.

The format of the time summary is $Uu\ Ss\ E\ P\% \ M$, where:

$U \Rightarrow$ User time in seconds

$S \Rightarrow$ System time in seconds

$E \Rightarrow$ Elapsed time, minutes:seconds

$P \Rightarrow$ Percentage of CPU time $((U + S)/E) * 100$

$M \Rightarrow$ Median text, data, and stack size use

q Terminate current input stream. If this is typed at top level, the simulator will exit back to the system; otherwise, input reverts to the previous input stream.

exit [n]

Exit to system, n is the reported status (default: 0).

Simulator parameters are set with the following commands. With no arguments, each of the commands simply prints the current value of the parameter.

decay [n]

Set decay parameter to n ns. (default: 0). If non-zero, it tells the number of ns. it takes for charge on a node to decay to X. A value of 0 implies no decay at all. You cannot specify this parameters separately for each node, but this turns out not to be a problem. See "**report**" command.

display [-][*cmdfile*][*automatic*]

set/reset the display modes, which are

cmdfile commands executed from command files are displayed to user before executing. The default is *cmdfile* = OFF.

automatic print out current display list (see "**d**" command) after completion of "**s**" or "**c**" command. The default is *automatic* = ON.

Prefacing the previous commands with a "-" turns off that display option.

model [*name*] **Set simulation model to one of the following:**

switch

Model transistors as voltage controlled switches. This model uses interval logic levels, without accounting for transistor resistances, so circuits with fighting transistors may not be accurately modelled. Delays may not reflect the *true* speed of the circuit as well.

linear

Model transistors as a resistor in series with a voltage controlled switch. This model uses a single-time-constant computed from the resulting RC network and uses a two-time-constant model to analyze charge sharing and spikes.

The default is the **linear** model. You can change the simulation model at any time -- even with events pending -- as only new calculations are affected. Without arguments, this command prints the current model name.

report [*level*]

When level is nonzero, report all nodes which are set to X because of charge decay, regardless on whether they are being traced. Setting level to zero disables reporting, but not the decay itself (see "decay" command).

stepsize [n]

Specify duration of simulation step or clock phase. n is specified in ns. (nanoseconds). Floating point numbers with up to 1 digit past the decimal point are allowed. Further decimals are truncated (i.e. 10.299 == 10.2).

unitdelay [*n*]

When nonzero, force all transitions to take *n* ns. Setting the parameter to zero disables this feature. The resolution is the same as for the "**stepsize**" command.

stats Print event statistics, as follows:

changes = 26077

punts (cns) = 208 (34)

punts = 0.79%, cons_punted = 16.35%

nevents = 28012; evaluations = 27972

Where *changes* is the total number of transistions recorded, *punts* is the number of punted events, (*cns*) is the number of consecutive punted events (a punted event that punted another event). The penultimate line shows the percentage of punted events with respect to the total number of events, and the percentage of consecutive punted events with respect to the number of punted events. The last line shows the total number of events (nevents) and the number of net evaluations.

Incremental simulation commands:

Irsim supports incremental changes to the network and resimulation of the resulting network. This is done incrementally so that only the nodes affected by the changes, either directly or indirectly, are re-evaluated.

update *filename*

Read net-change tokens from the specified file. The following net-change commands are available:

add type gate source drain length width [area]

delete type gate source drain length width [area]

move type gate source drain length width [area] g s d

cap node value

N node metal-area poly-area diff-area diff-perim

M node M2A M2P MA MP PA PP DA DP PDA PDP

thresh node low high

Delay node tplt tphl

For a detailed dscription of this file see netchange(5). Note that this is an experimental interface and is likely to change in the future.

Note that this command doesn't resimulate the circuit so that it may leave the network in an inconsistent state. Usually this command will be followed by an **isim** command (see below), if that is not the case then it's up to the user to initilize the state of the circuit. This command exists only for historical reasons and will probably disappear in the future. It's use is discouraged.

isim [*filename*]

Read net-change tokens from the specified file (see netchange(5)) and incrementally resimulate the circuit up to the current simulation time (not supported yet).

ires *n* **The incremental algorithm keeps track of nodes deviating from their past behavior as recorded in the network history. During resimulation, a node is considered to deviate from its history if it's new state is found to be different within *n* ns of its previous state. This command allows for changing the incremental resolution. With no arguments, it will print the current resolution. The default resolution is 0 ns.**

SEE ALSO

presim(1) (now obsolete) rsim(1) irsim-analyzer(3) sim(5) netchange(5)

IRSIM Tutorial (version 2.1)

1.0 IRSIM the switch-level simulator

This is a simple tutorial on irsim. It is not meant to replace the irsim manual which contains much more information. Try [man irsim](#) for information on all the commands that irsim supports. irsim is a switch-level simulator, in the sense it models the circuit at the level of transistors.


2.0 Creating a .sim file

Typically we will first use Magic to layout a circuit and then extract the circuit description and analyze it with irsim. Hence we will not typically create **.sim** circuit description manually, as is described below.

Manually creating a circuit description (.sim file) for irsim can be useful in some cases. For example it permits a transistor level design to be tested quickly, before a layout is generated.

Irsim files are just text files that contain information about a circuit. They have the extension .sim. The first time you use irsim, you will probably have to create the .sim file yourself using your favorite text editor.

Below is a CMOS circuit to implement $\sim a \sim b + \sim c$ where $\sim a$ means the inverse of a .



The first step in creating a **.sim** file for this circuit is to label all the nodes. Label power *Vdd* and ground *Gnd*

Irsim is not case sensitive so *vdd* and *gnd* will work also. You can label the other nodes anyway you want. Use any method you wish. For the circuit above, the inputs are labeled *a*, *b*, *c* and the output node is labeled out. Other internal nodes are labeled *n1* and *n2*. When using irsim, it is helpful to have the labeled circuit schematic available so that you know the names of the nodes that you want to probe. The [tut_irsim_gate.sim](#) file for the above circuit is shown below.

```
|units: 100  tech: scmos
|
|type  gate      source  drain  length  width
|-----
p      a         vdd     n1     2        4
p      b         n1     out    2        4
p      c         vdd     out    2        4

n      a         out     n2     2        4
n      b         out     n2     2        4
n      c         n2     gnd    2        4
```

The first four lines are comments. Any line that begins with the vertical bar '|' is a comment. This first line says that the technology used is scalable cmos.

Transistors are specified as follows:

```
|type gate source drain length width
|-----
```

where type is **p** for pmos and **n** for nmos. The gate, source, and drain refer to the terminals of the transistor. This is why you need a labeled circuit schematic to create a **.sim** file by hand. It is important that you get the connections correct. The length and width of the transistor is both in microns.

The first transistor specified in the example above is the top-left pmos transistor in the corresponding schematic. Study the rest of the irsim file and the circuit. It is easy to see how the two match. You can also enter in resistors and capacitors. See the irsim manual for the proper format to create these elements. Considering just transistors, you are now ready to run irsim.

3.0 Starting IRSIM

At the unix prompt type

```
irsim tut_irsim_scmos2um.prm tut_irsim_gate.sim
```

where

[tut_irsim_gate.sim](#) is the simulation file and

[tut_irsim_scmos2um.prm](#) is the parameter file of a 2 micron process technology.

irsim will tell you how many transistors you have and then display the irsim prompt shown below.

```
IRSIM>
```

3.1 Running IRSIM

You can enter in commands interactively or you can run a script. Scripts are just the **.cmd** files and they

contains commands that are exactly what you would type in at the irsim prompt. (See section 3.3 below).

Here is an example irsim session. The **IRSIM>** prompt precedes commands that were entered interactively. Text preceded by the vertical bar **|** is output from irsim.

```
IRSIM> stepsize 50
```

The basic idea of irsim is that you tell it which nodes to pull high, which nodes to pull low, which nodes to tristate, and then you tell irsim to run the simulation for a certain period of time. This period of time is the stepsize. The above command tells irsim that the stepsize is 50ns. The default stepsize = 100 nanoseconds.

```
IRSIM> w out c b a
```

The command 'w' tells irsim to *watch* the following nodes. So the above command tells it to watch the nodes *out*, *c*, *b*, and *a*. Irsim displays the nodes in the reverse order in the above command. Therefore the output order will be *a b c out*. This is just a matter of personal preference though. Enter in the nodes in any order you like.

```
IRSIM> d
```

d *displays* all the nodes that are being watched. You can also enter in something like 'd a b' which tells irsim to only display nodes *a* and *b*

```
| a=X b=X c=X out=X
| time = 0.0ns
```

at time zero, the values of the nodes are all undefined

```
IRSIM> l a b c
```

The **l** command forces the nodes to a logic *low* value of 0 The above command sets nodes *a* *b* and *c* to logic 0

```
IRSIM> s
```

s tells irsim to *simulate* for a certain period of time previously defined by the stepsize command. The default value is 100 ns, but we set it to 50ns above.

```
| a=0 b=0 c=0 out=1
| time = 50.0ns
```

Irsim displays the values of the nodes after each step because of the previous **d** command. The current time is also displayed. Note that time = 50 ns. This is the current simulation time now.

```
IRSIM> h c
```

h sets the following nodes to a logic *high* value. Therefore the above command sets node *c* to logic 1.

```
IRSIM> s
| a=0 b=0 c=1 out=1
| time = 100.0ns
```

step again

```

IRSIM> h b
IRSIM> s
| a=0 b=1 c=1 out=0
| time = 150.0ns

IRSIM> path out
| critical path for last transition of out:
| b -> 1 @ 100.0ns , node was an input
| out -> 0 @ 100.1ns (0.1ns)

```

The path command shows the critical path for the last transition of a node. The output shows that an input node 'b' changed to logic 1 at time = 100 ns. The node 'out' then transitioned low at time = 100.1 ns. Therefore it took .1 ns to go from high to low for the given input changes.

You can look back at previous commands to verify that this is indeed what happened.

Sometimes you may want to find out what the worst case Tplh or Tphl is. The path command helps you find this number. You do this by setting the circuit to some state and then force inputs to change that will cause a transition at the output. Some intelligence is required to figure out what the worst state is and what combination of input changes will cause the slowest output transition.

If you have a long list of inputs, it can be tiresome to keep using the l and h commands to set the logic values. The 'vector' command lets you group inputs together so that you can set them all quickly

```

IRSIM> vector in a b c

```

The above command tells irsim to group the nodes a b and c into a vector called 'in'

```

IRSIM> set in 000

```

This command tells irsim to set the vector in to 000 Therefore, node a = 0, node b = 0, and node c = 0

The following commands demonstrate how you can create a truth table using the vector 'in'. Note that you can do this much faster this way than by using the commands l and h.

```

IRSIM> s
| a=0 b=0 c=0 out=1
| time = 200.0ns
IRSIM> set in 001
IRSIM> s
| a=0 b=0 c=1 out=1
| time = 250.0ns
IRSIM> set in 010
IRSIM> s
| a=0 b=1 c=0 out=1
| time = 300.0ns
IRSIM> set in 011
IRSIM> s
| a=0 b=1 c=1 out=0
| time = 350.0ns
IRSIM> set in 100
IRSIM> s
| a=1 b=0 c=0 out=1
| time = 400.0ns

```

```
IRSIM> set in 101
IRSIM> s
| a=1 b=0 c=1 out=0
| time = 450.0ns
IRSIM> set in 110
IRSIM> s
| a=1 b=1 c=0 out=1
| time = 500.0ns
IRSIM> set in 111
IRSIM> s
| a=1 b=1 c=1 out=0
| time = 550.0ns
```

irsim has a built in graphical logic analyzer which runs under X to let you view waveforms. The 'analyzer' command sets up the analyzer window

```
IRSIM> analyzer a b c out
```

This tells irsim to display the nodes a b c and out in the analyzer window

3.2 Using the analyzer window

Go ahead and experiment with the analyzer window. It's pretty easy to use. Here are some of the things that you can do.

3.2.1. Print the waveforms to a postscript file

Go to the 'print' menu and select 'file'. You will be asked for a filename. Hitting return selects the filename shown in (). Once you have the postscript file, you can print it out on one of the network printers using `lp -d spot file.ps`.

3.2.2. Setting the width

3.2.2a. You can set the width of the view from either the menus or the slider on the bottom. From the menus, select 'window' and then 'set width'. You will be asked to enter the width in time steps. Use the 'move to' option under the 'window' menu to move to the desired time.

You can also use the 'zoom' menu to zoom in and out.

3.2.2b. The best way to get a feel for the slider on the bottom is to play with it. Using the left mouse button expands the view or shortens the view to the left depending on where you click. Clicking using the right mouse button expands or shortens the view to the right depending on where you click. You can grab the slider with the middle button and move it around.

3.2.3. Changing the order of the nodes

If you don't like the order that the nodes are displayed, you can click inside the name with the left mouse button, and drag it to a new place. Play around with this to see how it works.

3.2.4. Finding the time of an event

You can click inside the analyzer window where the waveforms are to get a vertical line. The time where this vertical line is is displayed above. This is useful if you want to find out when a rising edge occurs. The resolution is best when you are zoomed in. You have seen the commands 'l' and 'h' to set nodes to logic 0 or logic 1. The 'x' command will effectively put a node into tristate. It does this by taking the node off the input list.

The following is an example of using the x tristate command `IRSIM> x bus`

puts the node bus into tristate. An application of this is simulating a register. Assume you have a one bit register that is connected to a line called bus and is operated as follows. To write a value in your register, you need to drive the value onto the bus and then set a control line write to high. To read the register, you first need to stop driving the bus. Then set a control line read to high. The irsim commands you need to issue are shown below.

```
IRSIM> l read write           // set read and write control lines low
IRSIM> h bus                  // let's write a 1 into the register
IRSIM> s                      // step
IRSIM> h write                // write into the register
IRSIM> s                      //
IRSIM> l write                // stop writing
IRSIM> s                      //
IRSIM> l bus                  // set bus to 0 so we can see the bus
                               // transition from 0 to 1 when we read
                               // the register
IRSIM> s                      //
IRSIM> x bus                  // stop driving the bus so we don't get
                               // bus contention when we read the
                               // register
IRSIM> s                      //
IRSIM> h read                 // read the register; you should see a
                               // 0 to 1 transition
```

Irsim lets you keep a log of your commands and the outputs. To start recording, type

```
IRSIM> logfile anyfilename
```

Anything that you see in irsim from now on will be into the file 'anyfilename'

When you're done recording, type the following. `IRSIM> logfile`

Quitting IRSIM `IRSIM> exit`

3.3 Automation through .cmd files

Any irsim interactive command is valid. Use your favorite text editor to create this file. An example is shown below for the previous circuit.

```
stepsize 50
analyzer a b c out
vector in a b c
set in 000
s
set in 001
```

```

s
set in 010
s
set in 011
s
set in 100
s
set in 101
s
set in 110
s
set in 111
s

```

This example runs through all the possible combinations of inputs. The waveforms will be displayed in an analyzer window. You can run a .cmd file from the command line or within irsim. If the above file is called example.cmd, then at the Unix shell prompt, type

```
irsim scmos2um.prm circuit.sim -example.cmd
```

this will run the command file on circuit.sim

IRSIM commands

stepsize [n]	set simulation stepsize to n ns
s [n]	simulate for n nanoseconds (default= stepsize)
d [node]	Display status of concerned nodes
w [-]node	Watch a new node [or -remove a node] from display
h node1 node2 ...	High: Set list of nodes continuously to logic 1
l node1 node2 ...	Low: Set list of nodes continuously to logic 0
u node1 node2 ...	Undefined: Set list of nodes to "X" (undefined)
x node1 node2 ...	Tristate: Stop setting the list of nodes
vector label node(s)	define bit vector
set label bits	Set the bits to label
ana node1 node2 ...	(or analyzer) display nodes in analyzer window
clear	Clear waveform viewer display
flush [time]	flush history upto time (default : now)
logfile file.log	Turn logfile on
logfile	Turn logging off
q	quit

IMPORTANT. When you set a node high or low using the h or l commands, the node keeps being set to high or low (no matter what the circuit is trying to do to the node!) until you use the x command to stop setting the node.

Vectors

Since nodes typically are grouped into vectors, it is usually easier to look at *N*-bit quantities as single vector entities. The following commands can be used to define vectors and display them.

`vector name node1 node2 ...` define a new vector called *name* consisting of the list of nodes

`d name` display a vector as an array of bits

`ana name` add vector *name* to the waveform viewer

`set name value` set the bits of vector *name* using the binary string *value*

To set a vector to a hexadecimal number, use:

```
set name %xhexstring
```

For instance, you can now say: `set va %xff` if *va* is an 8-bit vector instead of `set va 11111111`. The prefix `%x` says that what follows is a hex constant.

A useful shortcut to defining arrays as long vectors is to say:

```
IRSIM> vector name a.b[{31:0}]
```

Clock Definition

The standard clock definition is shown below:

```
IRSIM> vector CLOCK CLK _CLK
IRSIM> clock CLOCK 01 10
```

The first line defines `CLOCK` to be a vector of two signals: `CLK` and `_CLK`. These signals are defined as globals in the file. The second line states that a single cycle of the clock consists of setting the vector first to `01` and then to `10`. Once this clock is defined, you can run the simulation for a clock step by saying:

```
IRSIM> c
```

The following runs the simulation for 10 cycles:

```
IRSIM> c 10
```

Acknowledgements

Original version by Williams, 11-apr-95. Downloaded from The IRSIM Tutorial version 2.0 Downloaded from www-leland.stanford.edu/class/ee272/doc/faq/irsim/ Modified by Fred DePiero at CalPoly 10/21/97. Also, EECS 314 Instructor Rajit Manohar at Cornell University. Finally, Modified by Francis G. Wolff at CWRU 7/18/00.

HOW TO DESIGN SIMULATABLE CMOS INTEGRATED CIRCUITS


William B. Ackerman
2 March 1985

There are many aspects of the successful design of custom VLSI circuits: high performance, reliability, testability, and so on. This memo will address the issue of maximizing the probability that the fabricated devices will actually work. Issues of performance, design discipline, etc., are covered elsewhere in VLSI design courses. In other words, this memo is intended to help you sleep better during the four months or so that your chip is being fabricated.

We assume that the chips are being prepared for the MOSIS CMOS fabrication, using the design tools commonly used by the MIT VLSI research group. In particular, we assume that the simulator "RNL" or "RSIM" is being used. We will use the term "RNL" to mean either RNL or RSIM. The underlying simulator in both programs is the same.

There are two important aspects of reliable design: avoiding circuit structures whose functioning in the fabricated device might differ from results of a simulation (hereafter referred to as "dangerous" circuit structures), and performing a successful simulation. That is, we want the simulator to say that the chip works, and we want to be confident that the simulator told the truth. As will be seen, these two considerations have much in common.


We will refer several times to this circuit, known as the "tiny XOR":




This is sometimes known as the "6 transistor XOR," because of the two other transistors that are presumably required to compute \bar{Q} from Q . The tiny XOR, or a similar circuit, is usually the chief culprit in unsuccessful simulations with RNL.

AVOIDING DANGEROUS CIRCUIT STRUCTURES

First, observe that standard "classical" CMOS gates are always completely safe. This is true whether they are simple:




or complex:




Any such gate will always work. Even if outrageously unbalanced transistor sizes are used, the circuit will always work, and RNL will successfully simulate it whether in switch or resistor mode (more about that later.) Therefore, use classical gates wherever possible.

The commonest non-standard circuit is the "pass gate":



It is common for two or more of these to be connected to some common "bus":




such that the enabling signals P_i are disjoint. The circuit may or may not turn off all of the P_i and expect the

previous voltage on the bus to remain in place due to charge storage on the bus's capacitance.


DYNAMIC CHARGE STORAGE

If charge is stored on the bus while all of the pass transistors are off, the first thing to be careful about is the nature of the bus's capacitance. The safest type of capacitor is the gate of a transistor, with as little stray capacitance as possible between the stored-charge node and other parts of the circuit. That is, the stored-charge node should be localized, not spread into the far reaches of the chip.




Remember that we want the capacitance to be to ground, not to other parts of the circuit that might have changing signals. So we want the capacitance of the inverter gate to be very much larger than the stray capacitances. One type of unavoidable stray capacitance is the gate to drain capacitance of the pass transistors themselves. When a pass transistor pair turns off, its two gates undergo large voltage swings. Fortunately, the voltage swings are in opposite directions, so their effects tend to cancel. Also, the gate to drain capacitance is much lower per unit area than the gate to substrate capacitance of the inverter. So unless the pass transistors are grossly bigger than those in the inverter, there should be no problem. When in doubt, simulate carefully with SPICE.

The foregoing applies even if the "bus" is driven by only one pass gate, as is common for dynamic latches:



This is also a common output circuit for precharged PLA's. The final inverter is important -- don't remove it.

One thing you must never do¹ is have a stored charge node drive a pass transistor:




If all of the P_i are off, so charge is stored on the bus, and E is turned on, the charge left on node G from the last time will combine with the charge on the bus, with very unsatisfactory results.

In fact, CNODE and RNL will probably not allow you to simulate such a circuit. CNODE only computes gate capacitance -- it does not recognize that sources, drains, and wiring have capacitance to the substrate, so RNL will refuse to believe that charge can be stored on such a node.


1. without giving the situation *very* serious thought

CHARGE SHARING AND FLIP-FLOPS

The other dangerous circuit structure related to pass transistors is a static latch.




When P is turned on, charge from the bus goes back and fights with the flip-flop. An unbuffered flip-flop such as the one shown can lose information if its outputs are interfered with. RNL may or may not show the malfunction -- it depends on RNL's perception of the capacitances, transistor ratios, thresholds, etc. The fabricated chip may or may not malfunction also -- it depends on the same things. Since CNODE and RNL use a very crude model of these parameters, don't take risks. This circuit is dangerous. Use one of these instead, to isolate the flip-flop from the effects of the pass transistor.



The way to think about all this is that transistor gates are very good listeners and never talk back. On the other hand, when you put a voltage on the source or drain of a pass transistor, it will fight with you when the gate is turned on, so you have to take steps to be sure that you win the fight.


CHARGE SHARING IN INCOMPLETE GATES

One can sometimes build a sort of nor gate with only three transistors:



It is easy to see that this is not a real NOR gate. Its defect is that if A and B are both low, so that the output is high, and A rises, the output will not fall. There are applications, however, in which this is not a problem, such as drivers for transistors pulling down precharged lines, such as in a PLA. B is a precharge signal that holds the output low. In such an application, the output's inability to fall when A rises while B is low is not a problem.


Assume that we are using this circuit in such a situation, so that its "defect" is not a problem. That is, our use of the circuit is *logically* correct. (If it is used in a manner that is not logically correct, RNL will say so.) There is a potential for a charge sharing bug. Suppose A and B are both high, though A had been low previously.



Because A had been low, it pulled point X up to VDD, even though it is not pulling it up now.

When B falls, we expect the output to stay low. However, the charge on X is placed on the output, making it rise somewhat. This can lead to a malfunction. *CNODE* and *RNL* will not show this malfunction, since they don't recognize that node X has capacitance. *EXCL* and *SPICE* will show it, and it can be observed in fabricated devices.

The solution to the problem is to exchange the order of the pullup transistors:




Now when B falls while A is high, the charge on X can't pass through to the output. *SPICE* simulations show that the isolation of X from the output is excellent.

THE TINY XOR


The sad fact is that RNL is essentially unable to simulate the tiny XOR. (Other simulators, such as *MOSSIM*, may have the same problem. *SPICE* does not.) Under certain circumstances, RNL has been observed to simulate it correctly: RNL was run in resistor mode rather than switch mode, and the transistors driving the inputs of the XOR were artificially broadened by a factor of 20. Even under these conditions RNL's behavior is problematical and uncertain. It has been known to fail even when the XOR was driven by transistors 100 times more powerful than the transistors in the XOR itself. Therefore, don't extract and simulate a 10000 transistor chip expecting this to be your lucky day. *Avoid simulating the tiny XOR and similar circuits.*

The problem is that, until RNL has solved the whole circuit, it doesn't know the value of the output. Until it knows whether P is being fought with, it doesn't know the value of P. It therefore doesn't know whether transistor B is on, causing the output to fight with Q. Therefore it doesn't know the value of Q, so it doesn't know whether transistor A is on, causing the output to fight with P. And so it goes. Don't argue with me. I'm just telling you what happens. Don't complain to the author or maintainer of RNL. Logic simulators are extremely complex programs, and the problem would be incredibly difficult to fix.



Despite RNL's disdain for the tiny XOR, there is actually nothing wrong with it, as long as you keep in mind that, like a pass gate, it is non-restoring (don't cascade lots of them without occasional restoring gates) and has a tendency to fight with its inputs (be careful about driving it from unbuffered flip-flops.) It is fast (about 2 ns with minimum geometries driving minimal loads), and compact.

It is perfectly reasonable to fabricate a chip containing these XOR circuits, *but you must never attempt to simulate such a circuit*. Therefore, when contemplating the use of such a circuit, first consider whether it is really important to do so. Where possible, compute XOR by conventional means. If speed or space requirements mandate the tiny XOR, you must do the following: Make two interchangeable instances in HPEDIT, one with the tiny XOR for fabrication and design rule checking, the other with an equivalent "safe" XOR for extraction. A "safe" equivalent XOR is this:



This solves RNL's problem: The circuit never fights with Q, so RNL will be able to figure out the value of Q. Having done that, it can resolve the fighting that occurs at P.

The safe XOR is more complex than a tiny XOR: it has two more transistors, and it needs access to power and ground. If space was a consideration in the decision to use the tiny XOR, this can be a problem. The solution is to violate the design rules in the "safe" XOR. The "safe" XOR is used only for node extraction, and the node extractor isn't bothered by half lambda gates or contact cuts. The simulation should also be independent of transistor sizes.


You must be extremely careful to check that the real XOR that will be fabricated will be logically equivalent, in the circuit, to the XOR that is simulated, since it will be the one part of the chip that will not be checked by computer.

Now reconsider whether you *really* want to use the tiny XOR.

OTHER CIRCUITS WITH TINY-XOR-LIKE PROBLEMS


The tiny XOR is not the only circuit exhibiting intractable behavior and requiring the extreme measures described above -- it is only the smallest such. Any circuit using pass transistors is suspect. Carry circuits in adders are a common source of trouble. This is why classical gates should be used wherever possible.

Here is an example of how insidious things can become. It arises in a full adder.




We have a 2-input multiplexer that, depending on the output of the LOGIC, passes either X or Y to point Z.

Until LOGIC has been solved, we don't know which of the pass gates is on and which is off. The simulator puts both into the unknown state, which causes X and Y to fight with each other. This may cause both X and Y to go into the unknown state. Since X is used as an input to LOGIC, the situation will never be resolved. Note that, if Y were absent, the problem would not arise.



X is only fighting with Z. Assuming Z is just a gate, it won't fight back. Hence X will win even while the state of the pass gate is not known.

How can we get out of this dilemma? Put an isolating gate before the pass gate.



(Two inverters are shown so that the logical function will be preserved. If one inverter can be used and its effect compensated for elsewhere, that's fine.)

Now the evil things happening in the pass gate will not affect X at the point where LOGIC senses it, so LOGIC will be able to operate correctly and eventually get the pass gate into a known state.

We must put in the isolation gate to make it simulate. Must we put it in the actual chip? It depends on whether it is harmful for the circuit to fight with inputs X and Y -- for example, whether these signals are coming from unbuffered flip-flops. Like the tiny XOR, the decision is yours.

SIMULATION HINTS

First, always use RNL or RSIM, version 3.3 or later. Do not use any "enhanced", "improved", "human engineered", or otherwise butchered versions. Do not use anything called "NL" or "CNL", even if someone tells you that they are "easier to use" or that "RNL doesn't do CMOS."

Second, it is almost always better to run RNL in "switch" mode rather than "resistor" mode. In switch mode RNL is completely unable to solve the voltage at a point that is being pulled in two directions by two transistors that are both conducting. Hence ratioed logic is impossible. However, ratioed logic is dangerous (fabrication parameters are hard to predict in MOSIS chips) and should be avoided in CMOS.

To use switch mode, put the line

```
(setq switch-level t)
```

into your simulation file, *after* the line

```
(load "n1.1")
```

The principal advantage of switch mode is that RNL's behavior in the presence of the tiny XOR and similar circuits is predictable and repeatable. If you use resistor mode to simulate a chip with such circuits, it might work most of the time. You might think that you have beaten the curse of the tiny XOR, only to find, when you do your final simulation, that RNL is taking many hours of computer time and getting nowhere. When run in switch mode, RNL's punishment is swift and decisive. It simply will not simulate the tiny XOR, and will quickly tell you so, by reporting a value of "X" for the output. When you get an unexplained "X," think carefully about whether there are XOR-like circuits, such as the adder described previously, and then remove them.

Of course, it is sometimes necessary to run RNL in transistor mode. For example, a chip might use bleeder transistors to hold precharged lines, or it might use ratioed flip-flops. In this case, it would be a good idea to first simulate the control logic, or any other questionable parts, in switch mode, to verify that they contain no dangerous circuit structures. One can then simulate the entire chip in resistor mode, confident that RNL won't suddenly get upset by something.